label.switching: An **R** Package for Dealing with the Label Switching Problem in MCMC Outputs

Panagiotis Papastamoulis

University of Piraeus

Abstract

Label switching is a well-known and fundamental problem in Bayesian estimation of mixture or hidden Markov models. In case that the prior distribution of the model parameters is the same for all states, then both the likelihood and posterior distribution are invariant to permutations of the parameters. This property makes Markov chain Monte Carlo (MCMC) samples simulated from the posterior distribution non-identifiable. In this paper, the **label.switching** package is introduced. It contains one probabilistic and seven deterministic relabeling algorithms in order to post-process a given MCMC sample, provided by the user. Each method returns a set of permutations that can be used to reorder the MCMC output. Then, any parametric function of interest can be inferred using the reordered MCMC sample. A set of user-defined permutations is also accepted, allowing the researcher to benchmark new relabeling methods against the available ones.

Keywords: label switching, mixture models, hidden Markov, MCMC, R.

1. Introduction

Mixture and hidden Markov models are a powerful tool for modelling a wide range of phenomena and they have been extremely useful in many fields (McLachlan and Peel 2000; Frühwirth-Schnatter 2006). Such applications include the presence of unobserved heterogeneity in the studied population or the approximation of unknown distributions, after deciding the proper number of latent states (or components). The complex nature of such models can be simplified by decomposing the model into simpler structures using latent (unobserved) variables. Data augmentation (Tanner and Wong 1987) is a standard technique exploited both by the EM algorithm (Dempster, Laird, and Rubin 1977) as well as the Gibbs sampler (Gelfand and Smith 1990).

Under a Bayesian perspective, MCMC estimation of the posterior distribution of such models

is quite straightforward. Gibbs sampling enables the simulation of a Markov chain from the joint posterior distribution of model parameters and latent variables. Nevertheless, the likelihood of such models is invariant to permutations of the components' labels and this property gives rise to the label switching phenomenon (Redner and Walker 1984; Jasra, Holmes, and Stephens 2005). It is well known that the presence of the label switching phenomenon in an MCMC sample serves as a necessary condition for the convergence of the chain to the target distribution. On the other hand, the presence of the phenomenon complicates the posterior inference.

Early attempts to solve the label switching were focused on the use of suitable identifiability constraints (Diebolt and Robert 1994; Richardson and Green 1997; Frühwirth-Schnatter 2001). However, it is not always possible to find such constraints and in most cases a single identifiability constraint will not respect the geometry of the posterior distribution. For these reasons, a variety of different relabeling algorithms has been proposed.

The purpose of this study is to introduce the **label.switching** package for R (R Core Team 2015), available from the Comprehensive R Archive Network at `http://CRAN.R-project.org/package=label.switching`, which can be used in order to deal with the label switching problem using various algorithms that have been proposed to the related literature. More specifically, the **label.switching** package consists of eight relabeling methods: ordering constraints, the Kullback-Leibler based algorithm (Stephens 2000b), the pivotal reordering algorithm (Marin, Mengersen, and Robert 2005; Marin and Robert 2007), the default and iterative versions of ECR algorithm (Papastamoulis and Iliopoulos 2010; Rodriguez and Walker 2014; Papastamoulis and Iliopoulos 2013; Papastamoulis 2014), the probabilistic relabeling algorithm (Sperrin, Jaki, and Wit 2010) and the data-based algorithm (Rodriguez and Walker 2014).

In many instances, it is required to draw meaningful comparisons between different relabeling algorithms or to benchmark novel methods against the existing ones. Both issues are addressed to the **label.switching** package. At first, the output of each relabeling method is reported in a way that the resulting single best clusterings are comparable among them. Moreover, the user can provide alternative sets of permutations arising from any (consistent) relabeling procedure and directly compare them to the available methods.

The rest of the paper is organized as follows. A short introduction to mixture models is given in Section 2. Section 2.1 discusses the label switching phenomenon and some helpful notation is introduced. The relabeling algorithms contained in the **label.switching** package are briefly reviewed in Section 3. Section 4 gives an overview of the package and the most important functions are described in detail. The practical implementation of the package is illustrated in Section 5 using two real datasets from classic mixture and hidden Markov models, as well as two simulated datasets from mixtures of bivariate normal distributions. The paper concludes in Section 6.

2. Mixture models

Let $\boldsymbol{x} = (x_1, \ldots, x_n)$ denote a sample of n (possibly multivariate) observations. Assume that $\boldsymbol{z} = (z_1, \ldots, z_n)$ is an unobserved (latent) sequence of state variables, with $z_i \in \{1, \ldots, K\}$, where $K > 1$ denotes a known integer. Let f denote a member of a parametric family of distributions $f \in \mathcal{F}_\Theta := \{f(\cdot|\theta) : \theta \in \Theta\}$, with $\Theta \subseteq \mathbb{R}^d$, for some $d \geqslant 1$. Conditionally to z_i

and a vector of parameters $\boldsymbol{\theta} = (\theta_1, \ldots, \theta_K)$ the observations are distributed according to

$$x_i|(z_i = k, \theta_k) \sim f(\cdot|\theta_k), \quad k = 1, \ldots, K \tag{1}$$

for $i = 1, \ldots, n$.

Assume that z_i, $i = 1, \ldots, n$, are independent random variables following the multinomial distribution with weights $\boldsymbol{w} := (w_1, \ldots, w_K)$, that is,

$$P(z_i = k|w_k) = w_k, \quad k = 1, \ldots, K \tag{2}$$

independently for $i = 1, \ldots, n$ and $\boldsymbol{w} \in \boldsymbol{W} := \{w_k > 0, k = 1, \ldots, K-1 : \sum_{k=1}^{K-1} w_k < 1; w_K := 1 - \sum_{k=1}^{K-1} w_k\}$. The marginal distribution of x_i is a finite mixture of K distributions:

$$x_i|\boldsymbol{\theta}, \boldsymbol{w} \sim \sum_{k=1}^{K} w_k f(x_i|\theta_k). \tag{3}$$

The conditional probability for observation i to belong to component k can be expressed as

$$p_{ik} = \frac{w_k f(x_i|\theta_k)}{w_1 f(x_i|\theta_1) + \ldots + w_K f(x_i|\theta_K)}, \quad i = 1, \ldots, n; \quad k = 1 \ldots, K. \tag{4}$$

We will refer to Equation 4 with the term classification probabilities. The observed likelihood of the mixture is defined as:

$$L(\boldsymbol{\theta}, \boldsymbol{w}; \boldsymbol{x}) = \prod_{i=1}^{n} \sum_{k=1}^{K} w_k f(x_i|\theta_k). \tag{5}$$

The complete likelihood of the model is written as:

$$L_c(\boldsymbol{\theta}, \boldsymbol{w}; \boldsymbol{x}, \boldsymbol{z}) = \prod_{i=1}^{n} w_{z_i} f(x_i|\theta_{z_i}). \tag{6}$$

2.1. Label switching phenomenon

Let $\mathcal{T}_K$ denote the set of permutations of $\{1, \ldots, K\}$. For any $\boldsymbol{\tau} = (\tau_1, \ldots, \tau_K) \in \mathcal{T}_K$ define the corresponding permutation of the component specific parameters, weights and allocations as: $\boldsymbol{\tau\theta} := (\theta_{\tau_1}, \ldots, \theta_{\tau_K})$, $\boldsymbol{\tau w} := (w_{\tau_1}, \ldots, w_{\tau_K})$ and $\boldsymbol{\tau z} := (\tau_{z_1}, \ldots, \tau_{z_n})$, respectively.

Notice now that the likelihood of the mixture model is invariant for any permutation of the parameters, that is,

$$L(\tau\boldsymbol{\theta}, \tau\boldsymbol{w}; \boldsymbol{x}) = L(\boldsymbol{\theta}, \boldsymbol{w}; \boldsymbol{x}), \tag{7}$$

for all $\tau \in \mathcal{T}_K$, $\boldsymbol{\theta} \in \Theta$, $\boldsymbol{w} \in \boldsymbol{W}$. Let $p(\boldsymbol{\theta}, \boldsymbol{w})$ denote the prior distribution of component specific parameters and weights of the model. It will be assumed that this prior is permutation invariant as well, that is,

$$p(\tau\boldsymbol{\theta}, \tau\boldsymbol{w}) = p(\boldsymbol{\theta}, \boldsymbol{w}), \tag{8}$$

for all $\tau \in \mathcal{T}_K$, $\boldsymbol{\theta} \in \Theta$, $\boldsymbol{w} \in \boldsymbol{W}$. A typical choice (Richardson and Green 1997; Frühwirth-Schnatter 2001; Marin *et al.* 2005) for the prior of the component parameters is to assume that $\boldsymbol{\theta} \sim \prod_{k=1}^{K} p(\theta_k)$, independently for $k = 1, \ldots, K$, for a specific family of distributions

$p(\cdot)$ which is common to all states. A common prior assumption on the mixture weights is a non-informative Dirichlet distribution.

Let $p(\boldsymbol{\theta}, \boldsymbol{w}|\boldsymbol{x}) \propto L(\boldsymbol{\theta}, \boldsymbol{w}; \boldsymbol{x})p(\boldsymbol{\theta}, \boldsymbol{w})$ denote the posterior distribution of the mixture model parameters. From Equations 7 and 8 it follows that the same invariance property holds for the posterior distribution, that is, $p(\tau\boldsymbol{\theta}, \tau\boldsymbol{w}|\boldsymbol{x}) = p(\boldsymbol{\theta}, \boldsymbol{w}|\boldsymbol{x})$, for all $\tau \in \mathcal{T}_K$, $\boldsymbol{\theta} \in \Theta$, $\boldsymbol{w} \in \boldsymbol{W}$. This implies that all marginal densities of component specific parameters and weights are coinciding. Now, if a simulated output from any MCMC sampler has converged to the symmetric posterior distribution, the generated values will be switching between the symmetric high posterior density areas.

This behavior is known as the label switching phenomenon and makes the generated MCMC sample non-identifiable. Hence, it is not straightforward to draw inference for any parametric function that depends on the labels of the components. In order to derive meaningful estimates, all simulated parameters should be switched to one among the $K!$ symmetric areas of the posterior distribution. This is done by applying suitable permutations of the labels $\{1, \ldots, K\}$ to each MCMC draw.

3. Algorithms

In this section we will describe the relabeling algorithms that are available in **label.switching**. It consists of seven deterministic and one probabilistic relabeling method, as shown in Table 1. The third column describes the necessary input of the algorithms, while the input notation is described in detail in Table 2, where m denotes the number of retained MCMC iterations.

For practical purposes it will be convenient to arrange all component specific parameters ($\boldsymbol{\theta}$) and weights/transition probabilities ($\boldsymbol{w}$) in a $K \times J$ matrix $\boldsymbol{\xi}$. The number of columns (J) of the global parameter vector $\boldsymbol{\xi} := (\boldsymbol{\theta}, \boldsymbol{w})$ is equal to the number of different types of parameters of the model. For example, if Equation 3 corresponds to a univariate normal mixture model, then there are $J = 3$ different types: ξ_{kj} denotes the mean ($j = 1$), variance ($j = 2$) and weight ($j = 3$), respectively, for component $k = 1, \ldots, K$. In case of a bivariate normal mixture, there are $J = 6$ parameter types for each component: two parameters for the mean vector, two variances, one covariance and one weight.

In the sequel, we will assume that an augmented sample $(\boldsymbol{\xi}^{(t)}, \boldsymbol{z}^{(t)})$, $t = 1, \ldots, m$, has been generated by an MCMC algorithm. Moreover, let $p_{ik}^{(t)}$, $t = 1, \ldots, m$, $k = 1, \ldots, K$, $i = 1, \ldots, n$ denote the corresponding classification probabilities across the MCMC run.

Function	Method	Input
`aic`	Ordering constraints	`mcmc, constraint`
`dataBased`	Data-based	`z, x, K`
`ecr`	ECR (default)	`z, zpivot, K`
`ecr.iterative.1`	ECR (iterative vs. 1)	`z, K`
`ecr.iterative.2`	ECR (iterative vs. 2)	`z, K, p`
`pra`	PRA	`mcmc, pivot`
`stephens`	Stephens	`p`
`sjw`	Probabilistic	`mcmc, z, x, complete`

Table 1: Available relabeling algorithms in the **label.switching** package.

Object	Type	Dimension	Details
`z`	Array (integer)	$m \times n$	simulated allocation vectors
`x`	Array	$n \times d$	observed data
`zpivot`	Numeric (integer)	n	pivot allocation vector
`p`	Array (real)	$m \times n \times K$	classification probabilities
`mcmc`	Array (real)	$m \times K \times J$	simulated parameters
`pivot`	Array (real)	$K \times J$	pivot parameter
`complete`	Function	–	complete log-likelihood function

Table 2: Input notation for the relabeling algorithms.

3.1. Ordering constraints

Imposing an artificial identifiability constraint to the MCMC sample is the simplest approach to the label switching problem. In such a case, the simulated MCMC output is permuted according to the ordering of a specific parameter. However, this approach works well only in cases that the selected constraint is able to separate the symmetric posterior modes, which is rarely true.

Algorithm 1 (Ordering constraints)**.**

1. *Choose a specific parameter type ξ_s, $s = 1, \ldots, J$.*

2. *For $t = 1, \ldots, m$ find the permutation $\tau^{(t)} \in \mathcal{T}_K$ that $\xi_{\tau_1 j}^{(t)} < \ldots < \xi_{\tau_K j}^{(t)}$.*

3.2. Stephens' method

One of the first principled solutions to the label switching problem was proposed by Stephens (2000b). The idea behind Stephens' algorithm is to make the permuted MCMC draws agree on the $n \times K$ matrix of classification probabilities. For this purpose, the Kullback-Leibler divergence between an averaged matrix of classification probabilities across the MCMC run and the classification matrix in each MCMC iteration is minimized in an iterative fashion. In general, Stephens' algorithm is very efficient in terms of finding the correct relabeling, but its drawback is the need to store the $m \times n \times K$ matrix `p` of classification probabilities.

Algorithm 2 (Kullback-Leibler relabeling)**.**

1. *Choose m initial permutations $\tau^{(t)}$ $t = 1, \ldots, m$ (usually set to identity).*

2. *For $t = 1, \ldots, m$, $k = 1, \ldots, K$ calculate $q_{ik} := \frac{1}{m} \sum_{t=1}^{m} p_{i\tau_k}^{(t)}$.*

3. *For $t = 1, \ldots, m$ find a permutation $\tau^{(t)} \in \mathcal{T}_K$ that minimizes $\sum_{i=1}^{n} \sum_{k=1}^{K} p_{i\tau_k}^{(t)} \log \left(\frac{p_{i\tau_k}^{(t)}}{q_{ik}} \right)$.*

4. *If an improvement is made to $\sum_{t=1}^{m} \sum_{i=1}^{n} \sum_{k=1}^{K} p_{i\tau_k}^{(t)} \log \left(\frac{p_{i\tau_k}^{(t)}}{q_{ik}} \right)$ go to step 2, finish otherwise.*

3.3. Pivotal reordering algorithm

The Pivotal Reordering Algorithm (PRA), proposed by Marin *et al.* (2005); Marin and Robert (2007), is a very simple geometrically-based solution to the label switching. The idea is to permute all simulated MCMC samples of parameters so that they are maximizing their similarity to a pivot parameter vector, as the complete MAP estimate. This is done by selecting the permutation that minimizes the Euclidean distance between the pivot and the set of permuted parameter vectors in each MCMC iteration. In principle, this method is a data-driven way to apply an artificial identifiability constraint on the parameter space.

Algorithm 3 (Pivotal reordering)**.**

1. *Define a pivot parameter vector:* $\boldsymbol{\xi}^* = (\xi^*_{kj})$, $k = 1, \ldots, K$, $j = 1, \ldots, J$.

2. *For* $t = 1, \ldots, m$ *find a permutation* $\tau^{(t)} \in \mathcal{T}_K$ *that maximizes* $\sum_{j=1}^{J} \sum_{k=1}^{K} \xi^{(t)}_{\tau_k j} \xi^*_{kj}$.

Note here that maximizing the dot product $\tau \boldsymbol{\xi}^{(t)} \cdot \boldsymbol{\xi}^*$ in step 2 is equivalent to minimizing the Euclidean distance between $\tau \boldsymbol{\xi}^{(t)}$ and $\boldsymbol{\xi}^*$.

3.4. ECR algorithms

ECR algorithm was originally proposed by Papastamoulis and Iliopoulos (2010) and it is based on the idea that equivalent allocation vectors are mutually exclusive from the label switching solution. Two allocation vectors are called equivalent if the first one arises from the second by simply permuting its labels. ECR algorithm partitions the set of allocation vectors into equivalence classes and selects a representative from each class. Then, the permutation needed to be applied in a given MCMC iteration is determined by the one that reorders the corresponding allocations in order to become identical to the representative of its class.

In the default version of ECR algorithm (`ecr`), equivalence classes are determined using a pivot allocation vector `zpivot`. The pivot is selected by choosing a high-posterior density point, such as the complete or non-complete Maximum A Posteriori (MAP) estimate. Rodriguez and Walker (2014) tried to relax the dependence of ECR algorithm to a pivot and proposed two iterative versions (`ecr.iterative.1` and `ecr.iterative.2`). The first algorithm is using as input only the simulated allocation variables and is initialized by a pivot selected at random. Then, the standard version of ECR is repeated until a fixed pivot has been found. Nevertheless, it is not guaranteed that this procedure will lead to a "good" pivot. The second iterative ECR algorithm requires the knowledge of classification probabilities across the MCMC run and it could be described as an allocation vectors version of Stephens' algorithm. Of course the problem of storing the matrix `p` applies to this method as well. However, as it will be demonstrated in the applications, `ecr.iterative.2` is significantly faster than `stephens`.

Algorithm 4 (ECR: Default version)**.**

1. *Define a pivot allocation:* $\boldsymbol{z}^* = (z^*_1, \ldots, z^*_n)$.

2. *For* $t = 1, \ldots, m$ *find a permutation* $\tau^{(t)} \in \mathcal{T}_K$ *that maximizes* $\sum_{i=1}^{n} I(\tau z^{(t)}_i = z^*_i)$.

Algorithm 5 (ECR: Iterative version 1)**.**

1. *Choose m initial permutations $\tau^{(t)}$, $t = 1, \ldots, m$ (usually set to identity).*

2. *Update the pivot: $z_i^* = mode\{\tau z_i^{(t)}; t = 1, \ldots, m\}$, $i = 1, \ldots, n$.*

3. *For $t = 1, \ldots, m$ find a permutation $\tau^{(t)} \in \mathcal{T}_K$ that maximizes $\sum_{i=1}^n I(\tau z_i^{(t)} = z_i^*)$.*

4. *If an improvement is made to $\sum_{t=1}^m \sum_{i=1}^n I(\tau z_i^{(t)} = z_i^*)$ go to step 2, finish otherwise.*

Algorithm 6 (ECR algorithm: Iterative version 2)**.**

1. *Choose m initial permutations $\tau^{(t)}$, $t = 1, \ldots, m$ (usually set to identity).*

2. *Update the pivot: $z_i^* = argmax\{p_{i\tau_k}^{(t)}; t = 1, \ldots, m\}$, $i = 1, \ldots, n$.*

3. *For $t = 1, \ldots, m$ find a permutation $\tau^{(t)} \in \mathcal{T}_K$ that maximizes $\sum_{i=1}^n I(\tau z_i^{(t)} = z_i^*)$.*

4. *If an improvement is made to $\sum_{t=1}^m \sum_{i=1}^n I(\tau z_i^{(t)} = z_i^*)$ go to step 2, finish otherwise.*

3.5. Probabilistic relabeling algorithm

Another method provided by the package is the probabilistic relabeling algorithm `sjw` of Sperrin *et al.* (2010). Under this concept, the permutation for each MCMC draw is treated as missing data with associated uncertainty. Then, an EM-type algorithm computes the expected values of $K!$ permutation probabilities per MCMC iteration, given an estimate of the parameter values. In the maximization step, this estimate is updated using a weighted average of all permuted parameters. This method requires a large amount of user input: the generated MCMC sample of parameters and latent allocation variables, the observed data and a function that computes the complete log-likelihood. The algorithm is not efficient when the number of components grows large due to the computational overload.

Algorithm 7 (Probabilistic relabeling)**.**

1. *Initialize an estimate of the parameters $\widehat{\boldsymbol{\xi}} = (\widehat{\boldsymbol{w}}, \widehat{\boldsymbol{\theta}})$ and repeat steps 2 and 3 until a fixed point is reached.*

2. *E-Step: For $t = 1, \ldots, m$, compute permutation probabilities $g_\tau^{(t)} \propto L_c(\tau\widehat{\boldsymbol{\xi}}|\boldsymbol{z}^{(t)})$, $\tau \in \mathcal{T}_K$.*

3. *M-Step: Update parameter estimate $\widehat{\boldsymbol{\xi}} = \frac{1}{m} \sum_{t=1}^m \sum_{\tau \in \mathcal{T}_K} g_\tau^{(t)} \tau \boldsymbol{\xi}^{(t)}$.*

3.6. Data-based relabeling

The data-based method of (Rodriguez and Walker 2014) is a deterministic relabeling algorithm. At first, a set of cluster centers m_{kr} and dispersion parameters s_{kr} is estimated for $k = 1, \ldots, K$, $r = 1, \ldots, d$. Next, the optimal permutations are defined as the ones minimizing a k-means type loss-function between the cluster pivots and the observed data, based on the simulated allocations in each MCMC iteration.

Algorithm 8 (Data-based relabeling)**.**

  1. *Find estimates m_{kr} and s_{kr}, $k = 1, \ldots, K$, $r = 1, \ldots, d$.*

  2. *For $t = 1, \ldots, m$, find a permutation $\tau \in \mathcal{T}_K$ that minimizes*

$$\sum_{k=1}^{K} \sum_{\ell=1}^{K} I(z_i^{(t)} = \tau_\ell) \sum_{\{i : \tau z_i^{(t)} = \ell\}} \sum_{r=1}^{d} \left(\frac{x_{ir} - m_{kr}}{s_{kr}} \right)^2 .$$

The estimates in step 1 are solely based on the observed data $\boldsymbol{x}$ and the simulated allocation variables $\{\boldsymbol{z}^{(t)}, \ t = 1, \ldots, m\}$. For more details, the reader is referred to algorithm 5 of Rodriguez and Walker (2014).

Finally, it is mentioned that algorithms `stephens`, `ecr`, `ecr.iterative.1`, `ecr.iterative.2` and `dataBased` are optimized using the package **lpSolve** (Berkelaar and others 2015) for the solution of the assignment problem (Burkard, Dell'Amico, and Martello 2009). This is a key-property for any computationally-efficient label switching solving algorithm, because in any other case the computational overload explodes as the number of components increases due to the computation of $K!$ quantities. By transporting the original problem into equivalent integer programming ones, the computational overload is avoided. The reader is referred to Rodriguez and Walker (2014). On the other hand, for each MCMC iteration, the `pra` and `sjw` algorithms require the computation of $K!$ dot products and permutation probabilities, respectively, so they are not suggested for large values of K.

4. Implementation in R

All previously described relabeling algorithms are available as stand-alone functions in the **label.switching** package, as shown in Table 1. The input of each function is described in Table 2. Each one of them returns a list of permutations. The user can conveniently call any combination of these methods using the function `label.switching`, which serves as the main call function of the package. Moreover, a set of user-defined permutations can be also supplied which is useful for comparison purposes. In this section we will describe the general call of `label.switching` and explain the input arguments and output values in detail.

4.1. Structure of the main function

The general usage is

```
label.switching(method, zpivot, z, K, prapivot, p, complete, mcmc,
    sjwinit, data, constraint, groundTruth, thrECR, thrSTE, thrSJW,
    maxECR, maxSTE, maxSJW, userPerm)
```

and the details of the implementation are described in the sequel.

`method` the desired combination of the available methods. It can be any non-empty subset of: `"ECR"`, `"ECR-ITERATIVE-1"`, `"ECR-ITERATIVE-2"`, `"PRA"`, `"STEPHENS"`, `"SJW"`, `"AIC"`, `"DATA-BASED"`. Also available is the option `"USER-PERM"` which corresponds to a user-defined set of permutations `userPerm`.

zpivot Obligatory only when "ECR" has been selected. It is a user-specified set of $d \geqslant 1$ pivots and it should be defined as an $d \times n \times K$ array. Each pivot should correspond to a high posterior-density area. Then, method "ECR" will be applied d times.

z $m \times n$-dimensional array corresponding to the set of simulated allocation vectors $\boldsymbol{z}^{(t)}$, $t = 1, \ldots, m$, with $z_i^{(t)} \in \{1, \ldots, K\}$, for all $t = 1, \ldots, K$. It is required by: "ECR", "ECR-ITERATIVE-1", "ECR-ITERATIVE-2", "SJW" and "DATA-BASED".

K Positive integer (at least equal to 2) indicating the number of mixture components. It is required by: "ECR", "ECR-ITERATIVE-1" and "DATA-BASED". If missing, then it is set to $\max\{z_i^{(t)} : t = 1, \ldots, m; i = 1, \ldots, n\}$.

prapivot Obligatory only when "PRA" has been selected. It is a user-specified $K \times J$ array corresponding to a high posterior-density area for the parameters of the mixture.

p $m \times n \times K$ matrix of classification probabilities as defined in Equation 4. Required by methods "STEPHENS" and "ECR-ITERATIVE-2".

complete Complete log-likelihood function of the model. Required by method SJW. The input should be a $K \times J$ vector of parameters as well as an n-dimensional vector of allocations. The function should return a single value which corresponds to the complete log-likelihood as defined in Equation 6.

mcmc $m \times K \times J$ array of simulated parameters across the MCMC run. Required by methods "PRA", "SJW" and "AIC".

sjwinit An index on the set $\{1, \ldots, m\}$ pointing at the MCMC iteration whose parameters will initialize the sjw algorithm (optional).

data the observed data $\boldsymbol{x} = (x_1, \ldots, x_n)$. Required by "SJW" and "DATA-BASED" methods.

constraint An (optional) integer between 1 and J corresponding to the parameter that will be used to apply the Ordering Constraint. If constraint = "ALL", all J ordering constraints are applied. Default value: 1.

groundTruth Optional integer vector of n allocations, which are considered as the "true" allocations of the observations. The output of all methods will be relabelled in a way that the resulting single best clusterings maximize their similarity with the ground truth.

thrECR An (optional) positive threshold controlling the convergence criterion for ecr.iterative.1 and ecr.iterative.2. Default value: 10^{-6}.

thrSTE An (optional) positive threshold controlling the convergence criterion for stephens. Default value: 10^{-6}.

thrSJW An (optional) positive threshold controlling the convergence criterion for sjw. Default value: 10^{-6}.

maxECR An (optional) integer controlling the maximum number of iterations for ecr.iterative.1 and ecr.iterative.2. Default value: 100.

maxSTE An (optional) integer controlling the maximum number of iterations for `stephens`. Default value: 100.

maxSJW An (optional) integer controlling the maximum number of iterations for `sjw`. Default value: 100.

userPerm An (optional) list with S user-defined permutations ($S \geqslant 1$). It is required only if `"USER-PERM"` has been chosen in `method`. In this case, `userPerm[[i]]` is an $m \times K$ array of permutations, $i = 1, \ldots, S$.

Let f denote the number of selected relabeling algorithms. The following values are returned.

permutations A list of f permutation arrays: `permutations[[i]][t,]` corresponds to the permutation that must be applied to the parameters generated in the t-th MCMC iteration, according to method i, $i = 1, \ldots, f$; $t = 1, \ldots, m$.

clusters n-dimensional vector of best clustering of the observations for each method.

timings the CPU time for the reordering part of each method, that is, the time to find the optimal permutations without taking into account the time spent by the user in order to compute the necessary input.

similarity $f' \times f'$ similarity matrix between the label switching solving methods in terms of their matching best-clustering allocations, where $f' = f$ if `groundTruth` is not supplied and $f' = f + 1$ in the opposite case.

The output of the `label.switching` function is reported in a way that all relabeling methods maximize the similarity of the estimated single best clusterings with respect to a reference allocation vector. For this purpose, the number of matching allocations between two vectors is used. This makes easier the comparison between the different methods. By default, the reference allocation vector corresponds to the estimated single best clustering according to the first algorithm provided in `method`. In case that `groundTruth` is supplied by the user, the reference allocation is set to the true one which is quite helpful in simulation studies.

It is evident that each algorithm requires different types of input. Methods `aic`, `dataBased` and `ecr-iterative-1` require only quantities that are directly available from the raw MCMC output and/or the observed data. The algorithms `ecr`, `ecr-iterative-2`, `pra` and `stephens` demand a few extra lines of coding that mainly handle quantities that are already in use while the MCMC sampler is running. Finally, `sjw` is more demanding as the user has to provide a function along with the MCMC output.

The supplementary function `permute.mcmc` reorders the MCMC sample (as stored in `mcmc`) according to the permutations returned by `label.switching`.

Usage:

```
permute.mcmc(mcmc, permutations)
```

Arguments:

mcmc $m \times K \times J$ array containing an MCMC sample.

`permutations` $m \times K$ array of permutations.

Value:

`output` reordered `mcmc` according to `permutations`.

5. Examples

5.1. Mixture of normal distributions: Fishery data

The fishery data is taken from Titterington, Smith, and Makov (1985) and it consists of $n = 256$ snapper length measurements. The histogram of the data is shown in Figure 1 (left) and it is obvious that the length of a randomly sampled fish exhibits strong heterogeneity. This is due to the fact that the age of each fish has not been recorded. The data has been previously analyzed as a mixture of K normal distributions, that is,

$$x_i \sim \sum_{k=1}^{K} w_k \mathcal{N}(\mu_k, \sigma_k^2),$$

independent for $i = 1, \ldots, n$. According to Frühwirth-Schnatter (2006), the number of components ranges from 3 to 5 and there are clearly four separated clusters in the MCMC draws. Here, we will consider a more challenging scenario with $K = 5$ components. The MCMC sampler described in package **bayesmix** (Grün 2015) is applied in order to simulate an MCMC sample of $m = 10000$ iterations from the posterior distribution of $\{z, \mu, \sigma^2, w\}$, following a burn-in period of 1000. This is done with the following commands.

```
R> library("bayesmix")
R> data("fish", package = "bayesmix")
```


Figure 1: Histogram of the fishery data and time series of the fetal lamb movements.
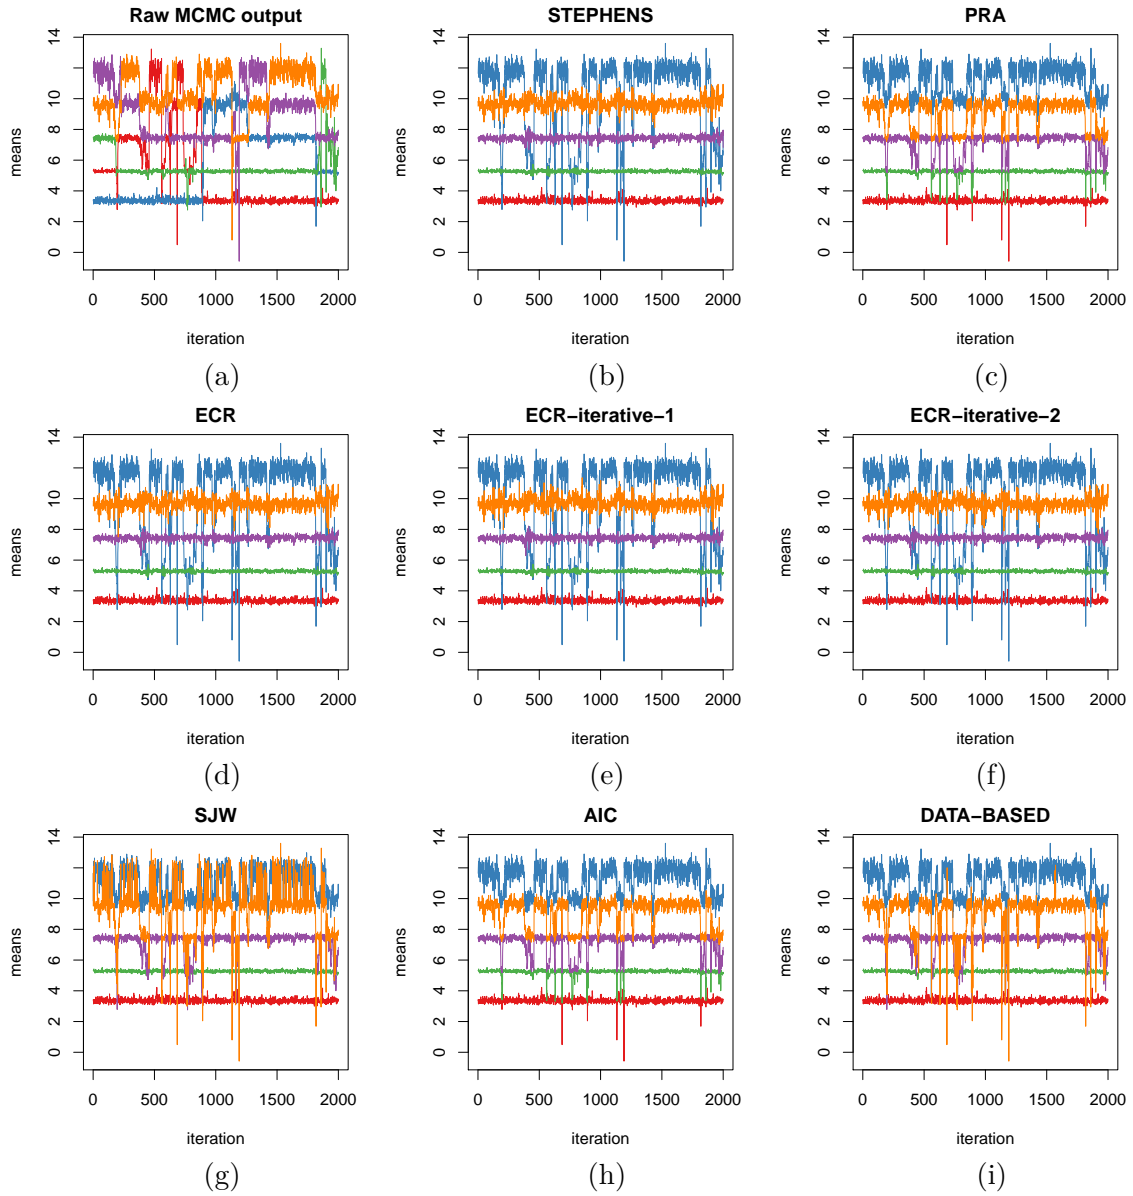
Figure 2: Fishery data. (a): Raw MCMC sample for μ_k. (b), (c), (d), (e), (f), (g), (h), (i): Reordered MCMC sample by applying the permutations returned by `label.switching` function, according to methods: Stephens, PRA, ECR, ECR-iterative-1, ECR-iterative-2, SJW, AIC and Data-based, respectively.

```
R> x <- fish[ , 1]
R> n <- length(x)
R> K <- 5
R> m <- 11000
R> burn <- 1000
R> model <- BMMmodel(fish, k = K, initialValues = list(S0 = 2),
+    priors = list(kind = "independence", parameter = "priorsFish",
+    hierarchical = "tau"))
```

```
R> control <- JAGScontrol(variables = c("mu", "tau", "eta", "S"),
+    burn.in = burn, n.iter = m, seed = 10)
R> mcmc <- JAGSrun(fish, model = model, control = control)
```

The raw MCMC output for μ_k, $k = 1, \ldots, 5$ is shown in Figure 2(a) (every 5th iteration displayed). It is obvious that the label switching phenomenon has occurred. Note that a simple ordering constraint to the means is not able to successfully isolate one of the symmetric high posterior-density areas. Next, we will apply the function `label.switching` considering all the presented relabeling algorithms. In order to do this, we have to compute all the related information that is required as input for each method. At first, the MCMC output is converted into an $m \times K \times J$ array (`mcmc.pars`), where $J = 3$ denotes the number of different parameter types for the normal mixture model: means (`mcmc.pars[ , , 1]`), variances (`mcmc.pars[ , , 2]`) and weights (`mcmc.pars[ , , 3]`). Finally, the generated allocation variables are stored to $m \times n$ array `z`:

```
R> J <- 3
R> mcmc.pars <- array(data = NA, dim = c(m, K, J))
R> mcmc.pars[ , , 1] <- mcmc$results[-(1:burn), (n+K+1):(n+2*K)]
R> mcmc.pars[ , , 2] <- mcmc$results[-(1:burn), (n+2*K+1):(n+3*K)]
R> mcmc.pars[ , , 3] <- mcmc$results[-(1:burn), (n+1):(n+K)]
R> z <- mcmc$results[-(1:burn), 1:n]
```

Stephens' method as well as the second iterative version of ECR algorithm need the $m \times n \times K$ array of component membership probabilities p_{ik}, as defined in Equation 4, for each MCMC iteration. These probabilities are stored to array `p` as follows.

```
R> p <- array(data = NA, dim = c(m, n, K))
R> for (iter in 1:m) {
+    for(i in 1:n) {
+      kdist <- mcmc.pars[iter, , 3]*dnorm(x[i], mcmc.pars[iter, , 1],
+      sqrt(mcmc.pars[iter, , 2]))
+      skdist <- sum(kdist)
+      for(j in 1:K) {
+        p[iter, i, j] = kdist[j]/skdist
+      }
+    }
+ }
```

Method `sjw` demands to provide as input a function that computes the complete log-likelihood of the classic mixture model, as defined by taking the logarithm of Equation 6. The next code accepts as input a dataset of univariate observations (`x`), an n-dimensional integer vector of allocations (`z`) and a $K \times J$ array of mixture parameters (means, variances, weights).

```
R> complete.normal.loglikelihood <- function(x, z, pars) {
+    g <- dim(pars)[1]
+    n <- length(x)
+    logl <- rep(0, n)
+    logpi <- log(pars[ , 3])
```

```
+     mean <- pars[ , 1]
+     sigma <- sqrt(pars[ , 2])
+     logl <- logpi[z] + dnorm(x, mean = mean[z], sd = sigma[z], log = TRUE)
+     return(sum(logl))
+   }
```

The function `complete.normal.loglikelihood` will be also used for the determination of an MCMC iteration that corresponds to a high density area. Next, the allocation and parameters of this iteration will be used as pivot by the functions `ecr` and `pra`, respectively. After evaluating the complete log-likelihood function for the 10000 MCMC iterations, we obtained that the maximum value corresponds to iteration `mapindex = 4839`.

We will also use an ordering constraint to the simulated means. Since this parameter type corresponds to `mcmc.pars[ , , j]` for `j = 1`, we should use `constraint = 1`. Now, we can apply the available algorithms using the following command.

```
R> library("label.switching")
R> set <- c("STEPHENS", "PRA", "ECR", "ECR-ITERATIVE-1", "ECR-ITERATIVE-2",
+     "SJW", "AIC", "DATA-BASED")
R> ls <- label.switching(method = set, zpivot = z[mapindex, ], z = z,
+     K = K, prapivot = mcmc.pars[mapindex, , ], p = p, constraint = 1,
+     sjwinit = mapindex, complete = complete.normal.loglikelihood,
+     mcmc = mcmc.pars, data = x)
R> ls$timings
```

The last command returns the CPU time per method, which is shown in first line of Table 3. The MCMC draws can be reordered applying the permutations contained in `ls$permutations` using the function `permute.mcmc`. Figure 2(b)-(h) contain the reordered output for the simulated values of μ_k, $k = 1, \ldots, 5$ (every 5-th iteration is displayed). We conclude that the results of methods `ecr`, `ecr-iterative-1`, `ecr-iterative-2` and `stephens` are quite similar to each other. The reordered values indicate that the component with the largest mean (blue colored trace) exhibits a multimodal behavior: there is a main mode at 12 and a minor one between 5 and 6. On the other hand, `pra` and `dataBased` algorithms are driven by the generated values of the means and the resulting reordering is quite similar to the one that results from an ordering constraint to μ_k. A similar behavior is observed from `sjw` algorithm.

Finally, the function `label.switching` returns the single best clusterings of the n observations among the K groups. This is simply done by calculating the mode of the reordered allocation vectors z_i, $i = 1, \ldots, n$. The proportion of the matching allocations between any pair of the available methods is returned by `ls$similarity` and it is shown in the lower diagonal of Table 4.

Data	K	n	m	steph	pra	ecr	ecr-1	ecr-2	sjw	aic	d-based
Fish	5	256	10000	56.9	3.5	3.4	17.8	13.4	613.7	0.1	11.4
Lamb	4	240	10000	61.0	0.8	2.7	11.4	11.0	598.5	0.1	8.5
Multivariate 1	4	100	10000	25.7	0.8	2.5	10.7	6.0	695.9	0.1	11.0
Multivariate 2	9	280	15000	431.3	NA	9.1	58.1	50.2	NA	0.1	61.1

Table 3: CPU times in seconds per relabeling method.

	steph	pra	ecr	ecr-1	ecr-2	sjw	aic	d-based
stephens		0.950	0.992	0.996	0.983	0.950	0.879	0.842
pra	0.996		0.958	0.954	0.938	0.904	0.833	0.863
ecr	1.000	0.996		0.996	0.975	0.942	0.871	0.850
ecr-iter-1	1.000	0.996	1.000		0.979	0.946	0.875	0.846
ecr-iter-2	1.000	0.996	1.000	1.000		0.966	0.899	0.825
sjw	0.992	0.996	0.992	0.992	0.992		0.896	0.792
aic	0.996	1.000	0.996	0.996	0.996	0.996		0.720
d-based	0.996	1.000	0.996	0.996	0.996	0.996	1.000	

Table 4: `ls$similarity`: Proportion of matching allocations for the single best-clusterings for each relabeling algorithm: Lower diagonal: Fish data. Upper diagonal: Fetal lamb data.

5.2. Poisson hidden Markov model: Fetal lamb data

A generalization of the classic mixture model set-up is to assume that the latent variables are forming an (unobserved) Markov chain. Let $\boldsymbol{w} = (w_{\ell k})$, $\ell, k \in \{1, \ldots, K\}$, denote an $K \times K$ matrix of transition probabilities. The conditional distribution of latent variables now is written as:

$$P(z_i = k | z_{i-1} = \ell, \boldsymbol{w}) = w_{\ell k}, \quad k = 1, \ldots, K. \tag{9}$$

The sequence $(z_1, \ldots, z_n)$ is unobserved and this justifies the term hidden Markov model. In this case, the complete likelihood is defined as:

$$L_c(\boldsymbol{\theta}, \boldsymbol{w} | \boldsymbol{x}, \boldsymbol{z}) = \pi_{z_1} f(x_1 | \theta_{z_1}) \prod_{i=2}^{n} w_{z_{i-1} z_i} f(x_i | \theta_{z_i}), \tag{10}$$

where π_k, $k = 1, \ldots, K$ denotes the left eigenvector of the transpose transition matrix $\boldsymbol{w}^T$, which corresponds to eigenvalue 1. Finally, the weights in Equations 3 and 4 are replaced by π_k, $k = 1, \ldots, K$. For an overview of hidden Markov model theory and applications the reader is referred to Cappé, Moulines, and Ryden (2005); Frühwirth-Schnatter (2006).

The fetal lamb data (Leroux and Putterman 1992) consists of $n = 240$ body movement measurements of a fetal lamb at consecutive 5 second intervals. The dependence of consecutive measurements is a sensible assumption here: a measurement at a specific intensity is quite likely to be followed by a similar one, as displayed in the corresponding time series in Figure 1 (right). Frühwirth-Schnatter (2001; 2006) modelled this time series as a Poisson process where the intensity changes according to a K-state hidden Markov process, that is,

$$
\begin{aligned}
x_i | z_i &\sim \mathcal{P}(\lambda_{z_i}), \text{ independent for } \quad i = 1, \ldots, n \\
P(z_i = k | z_{i-1} = \ell) &= w_{\ell k}, \quad k, \ell \in \{1, \ldots, K\}, \quad i = 2, \ldots, n,
\end{aligned}
$$

with $\sum_{k=1}^{K} w_{\ell k} = 1$ for all $\ell = 1, \ldots, K$, given some initial distribution $(\pi_1^{(0)}, \ldots, \pi_K^{(0)})$ for $i = 1$. Given K, the parameters to be estimated is the vector of intensities $\boldsymbol{\lambda} = (\lambda_1, \ldots, \lambda_K)$ and the matrix of state-transition probabilities $\boldsymbol{w} = (w_{\ell k})$, $\ell, k \in \{1, \ldots, K\}$. The number of states is estimated between 3 and 4. We will assume that $K = 4$ and run a Gibbs sampler for $m = 10000$ iterations after a burn-in period of 10000, using the same prior assumptions as discussed in Frühwirth-Schnatter (2001).
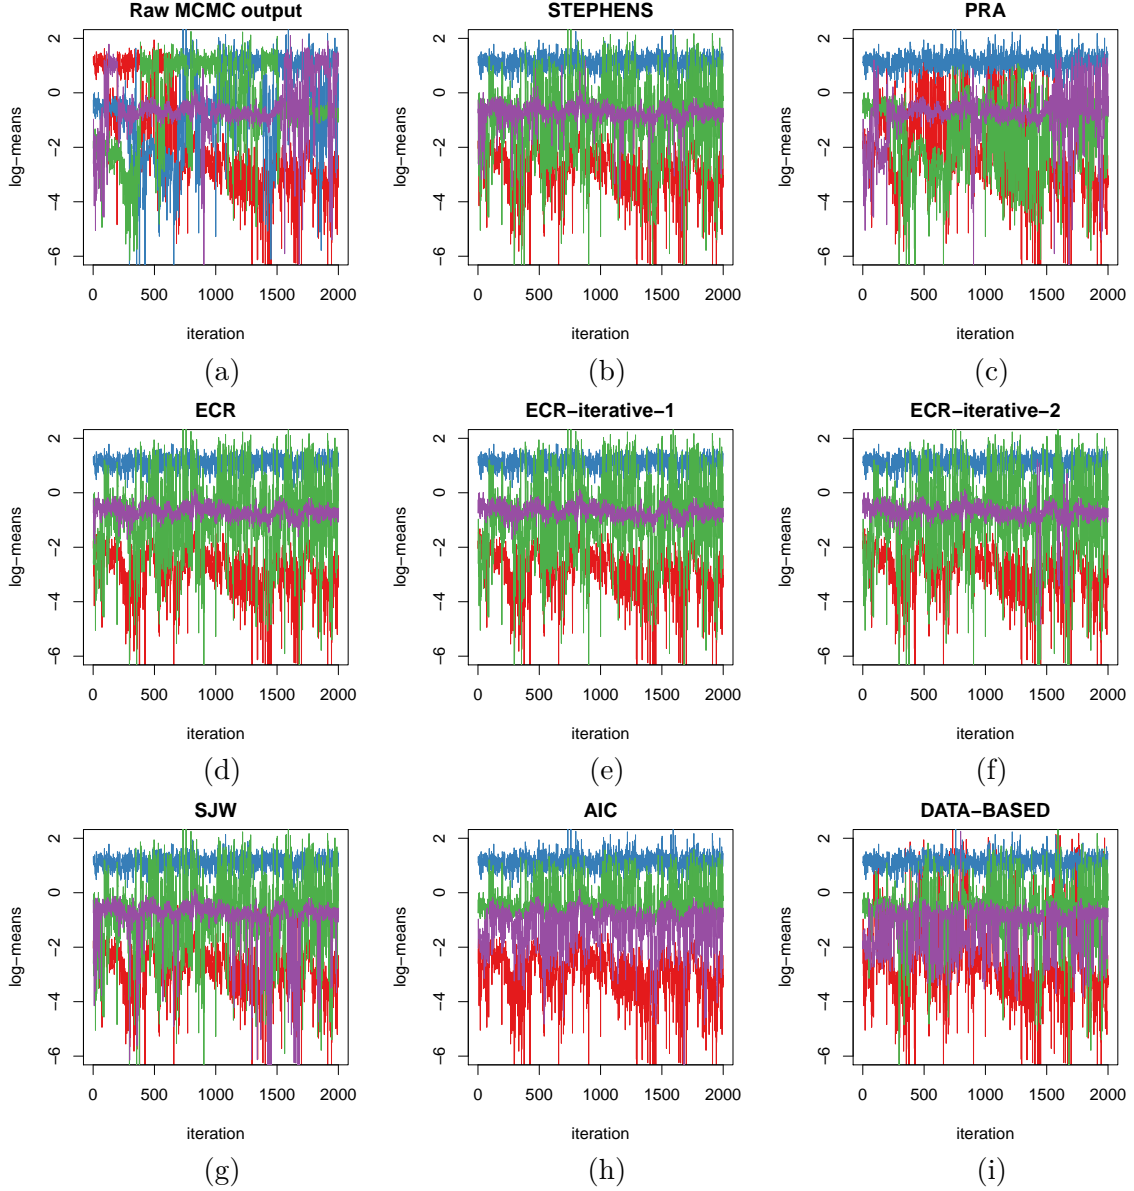
Figure 3: Lamb data ($K = 4$). (a): Raw MCMC sample of $\log \lambda_k$, $k = 1, \ldots, 4$. (b), (c), (d), (e), (f), (g), (h), (i): Reordered values by applying the permutations returned by `label.switching` function, according to methods: `stephens`, `pra`, `ecr`, `ecr-iter-1`, `ecr-iter-2`, `sjw`, `aic` and `dataBased` respectively.

At first we define the complete likelihood of the model, needed by `sjw` and `ecr` methods. The function `complete.hmm.poisson.loglikelihood` accepts as input the observed count data (`x`), a vector of allocations (`z`) and an $K \times J$ dimensional array (`pars`) of parameters with $J = K + 2$, where: $j = 1$ corresponds to Poisson means ($\boldsymbol{\lambda}$), $j = 2, \ldots, K + 1$ corresponds to K the columns of the state transition matrix $\boldsymbol{w}$ and $j = K + 2$ corresponds to the left eigenvectors of $\boldsymbol{w}$. Note here that given $\boldsymbol{w}$ it is not necessary to save the eigenvectors, but this will make the computation of the complete log-likelihood a little bit faster. The value returned corresponds to the logarithm of Equation 10.

```
R> complete.hmm.poisson.loglikelihood <- function(x, z, pars) {
+    post <- sum(dpois(x, pars[z, 1], log = TRUE))
+    logprobs <- log(pars[ , 2:(K+1)])
+    ev <- pars[ , K+2]
+    for(i in 2:n) {
+       post <- post + logprobs[z[i-1], z[i]]
+    }
+    post <- post + log(ev[z[1]])
+    return(post)
+ }
```

The full dataset is available in the **label.switching** package, while the MCMC sampler is provided as supplemental material. In the end, we have saved the MCMC output to an $m \times K \times J$ array (mcmc.pars). The raw MCMC output of $\log \lambda_k$, $k = 1, \ldots, 4$ is shown in Figure 3(a) (every 5th iteration displayed) where the label switching phenomenon is vividly illustrated. Next we apply the relabeling algorithms using the following command.

```
R> set <- c("STEPHENS", "PRA", "ECR", "ECR-ITERATIVE-1", "ECR-ITERATIVE-2",
+    "SJW", "AIC", "DATA-BASED")
R> ls <- label.switching(method = set, zpivot = z[mapindex, ], z = z, K = 4,
+    prapivot = mcmc.pars[mapindex, , ], p = p, mcmc = mcmc.pars, data = x,
+    complete = complete.hmm.poisson.loglikelihood, constraint = 1)
```

Note that for the default version of ECR and PRA algorithms we provided the pivot that correspond to iteration mapindex = 3258, that is, the allocation and parameters that correspond to the iteration that the maximum value of the complete likelihood was observed. After applying the function permute.mcmc using the resulting permutations, the reordered output of $\log \lambda_k$, $k = 1, \ldots, 4$ is shown in Figures 3(b)-(i). We conclude that almost all methods suggest that the values of the green-colored component (λ_3) have a very large variance compared to the rest. However, this is not the case for pra, aic, and to lesser extent for the dataBased algorithm, where the reordering does not seem to respect the posterior distribution topology. The proportion of matching allocations between the available methods is returned by ls$similarity and they are shown in the upper diagonal of Table 4. Finally, the CPU time per method is shown in second row of Table 3.

It is interesting to note here that most relabeling algorithms assign no observations (in terms of their single best clusterings) into the third (green) component. In particular, the estimated number of observations assigned to each cluster is:

```
R> frequency <- apply(ls$clusters, 1, function(y) {
+    freq <- numeric(K)
+    for(j in 1:K) {
+       freq[j] <- length(which(y == j))
+    }
+    return(freq)
+ })
+ rownames(frequency) <- 1:K
+ frequency
```

	steph	pra	ecr	ecr-1	ecr-2	sjw	aic	data	user	true
steph		1.000	1.000	1.000	1.000	1.000	0.880	1.000	1.000	0.940
pra	-		1.000	1.000	1.000	1.000	0.880	1.000	1.000	0.940
ecr	0.993	-		1.000	1.000	1.000	0.880	1.000	1.000	0.940
ecr-1	0.996	-	0.989		1.000	1.000	0.880	1.000	1.000	0.940
ecr-2	1.000	-	0.993	0.996		1.000	0.880	1.000	1.000	0.940
sjw	-	-	-	-	-		0.880	1.000	1.000	0.940
aic	0.896	-	0.889	0.889	0.896	-		1.000	0.880	0.830
data	1.000	-	0.993	0.996	1.000	-	0.896		1.000	0.940
user	0.971	-	0.971	0.968	0.971	-	0.892	0.971		0.940
true	0.929	-	0.929	0.929	0.929	-	0.836	0.929	0.918	

Table 5: `ls$similarity`: Proportion of matching allocations for the single best-clusterings between the relabeling algorithms, the user-defined permutations and the true allocations for the first and second multivariate dataset (upper and lower diagonal, respectively).

	STEPHENS	PRA	ECR	ECR-ITERATIVE-1	ECR-ITERATIVE-2	SJW	AIC	DATA-BASED
1	128	117	126	127	132	140	136	90
2	6	6	6	6	6	6	6	6
3	0	6	0	0	0	0	21	0
4	106	111	108	107	102	94	77	144

suggesting that the four-state hidden Markov model might be overparameterized for the fetal lamb dataset.

5.3. Multivariate normal mixtures

In this section, the **label.switching** package is applied to simulated data from mixtures of multivariate normal distributions. Let $\boldsymbol{x}_i \in \mathbb{R}^d$ denote a d-dimensional random vector and assume that

$$\boldsymbol{x}_i \sim \sum_{k=1}^{K} w_k \mathcal{N}_d(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k),$$

independent for $i = 1, \ldots, n$, with $\boldsymbol{\mu}_k \in \mathbb{R}^d$ and $\boldsymbol{\Sigma}_k \in \mathcal{M}^{d \times d}$, $k = 1, \ldots, K$, where $\mathcal{M}^{d \times d}$ denotes the space of $d \times d$ positive definite matrices. Let $\boldsymbol{\Lambda}_k := \boldsymbol{\Sigma}_k^{-1}$ and a priori assume a normal-Wishart prior distribution

$$\boldsymbol{\mu}_k, \boldsymbol{\Lambda}_k | \boldsymbol{\mu}_0, \beta, \boldsymbol{W}, \nu \sim \mathcal{N}(\boldsymbol{\mu}_k | \boldsymbol{\mu}_0, (\beta \boldsymbol{\Lambda}_k)^{-1}) \mathcal{W}(\boldsymbol{\Lambda}_k | \boldsymbol{W}, \nu),$$

independent for $k = 1, \ldots, K$, given the constant hyper-parameters $\beta > 0$, $\nu > d-1$, $\boldsymbol{\mu}_0 \in \mathbb{R}^d$ and $\boldsymbol{W} \in \mathcal{M}^{d \times d}$. The mixture weights are a priori distributed according to a non-informative Dirichlet distribution. The reader is referred to the supplementary material for the details of the hyper-parameters of the prior distributions.

We simulated two datasets of $n = 100$ and 280 observations from a bivariate ($d = 2$) mixture with $K = 4$ and 9 components, respectively. The real values used to generate the first dataset were chosen as $\boldsymbol{\mu}_k = 2.5 \left(\cos \frac{(k-1)\pi}{4}, \sin \frac{(k-1)\pi}{4} \right)^t$, $w_k = 1/K$, $\Sigma_{11k} = \Sigma_{22k} = 1$, $\Sigma_{12k} = \Sigma_{21k} = 0$, $k = 1, \ldots, 4$. The real values for the second dataset were chosen as

$\boldsymbol{\mu}_k = 6 \left(\cos \frac{(k-1)\pi}{8}, \sin \frac{(k-1)\pi}{8} \right)^t$, $\Sigma_{11k} = \Sigma_{22k} = 1$, $\Sigma_{12k} = \Sigma_{21k} = 0$, $w_k = 0.1$, for $k = 1, \ldots, 8$ and for the last component: $\mu_9 = (0,0)$, $\Sigma_{119} = \Sigma_{229} = 4$, $\Sigma_{129} = \Sigma_{219} = 0$ and $w_9 = 0.2$.

The Gibbs sampler was implemented next, using the package **mvtnorm** (Genz *et al.* 2016) in order to simulate from the full conditional distributions. The source code is provided in the supplementary material (`gibbsSampler` function). This functions returns the following objects: `mcmc`, `MLindex`, `z` and `p`, which correspond to the simulated parameters, the index which corresponds to the iteration where the maximum value of the complete likelihood is observed, the simulated allocations and the classification probabilities, respectively. More specifically, `mcmc` is an $m \times K \times J$ array, where $J = d + d(d+1)/2 + 1$ denotes the number of different parameter types for the bivariate normal mixture. Hence, `mcmc[t, k, 1]` $= \mu_{1k}^{(t)}$, `mcmc[t, k, 2]` $= \mu_{2k}^{(t)}$, `mcmc[t, k, 3]` $= \Sigma_{11}^{(t)}$, `mcmc[t, k, 4]` $= \Sigma_{22}^{(t)}$, `mcmc[t, k, 5]` $= \Sigma_{12}^{(t)}$ and `mcmc[t, k, 6]` $= w_k^{(t)}$, $k = 1, \ldots, K$, $t = 1, \ldots, m$. The function is called as follows.

```
R> gs <- gibbsSampler(iterations = iterations, K = K, x = x, burn = burn)
R> zChain <- gs$z
R> mcmc.pars <- gs$mcmc
R> pivot <- gs$MLindex
R> allocProbs <- gs$p
```

For the first dataset we set `K = 4`, `iterations = 11000` and `burn = 1000` and for the second dataset we set `K = 9`, `iterations = 20000` and `burn = 5000`.

All relabeling algorithms are applied to the first dataset, but `pra` and `sjw` are excluded to the second one due to the large number of permutations ($K!$) that should be computed for each MCMC iteration. For the `sjw` algorithm, the `complete.bivariate.normal.loglikelihood` function (available in the supplementary material) returns the complete log-likelihood function of the bivariate normal mixture. The ordering constraint will be applied to μ_{1k}, $k = 1, \ldots, K$. We will also provide an additional set of permutations using the `userPerm` option. Assume that after visual inspection of the MCMC draws, the user wishes to check whether the MCMC sample is identifiable by imposing an ordering constraint to $\mu_{1k} - 2\mu_{2k}$, $k = 1, \ldots, K$. This is easily done using the following commands.

```
R> newMCMC <- array(data = NA, dim = c(iterations, K, 7))
R> newMCMC[ , , 1:6] <- mcmc.pars
R> for(k in 1:K){
+        newMCMC[ , k, 7] <- mcmc.pars[ , k, 1] - 2*mcmc.pars[ , k, 2]}
R> newConstraint <- aic(newMCMC, constraint = 7)
```

Now apply all relabeling algorithms using the `label.switching` command, by parsing also the user-defined permutations in order to compare them to the rest of the methods as follows. For the first dataset

```
R> set <- c("STEPHENS", "PRA", "ECR", "ECR-ITERATIVE-1", "ECR-ITERATIVE-2",
+    "SJW", "AIC", "DATA-BASED", "USER-PERM")
R> ls <- label.switching(method = set, zpivot = zChain[pivot, ], z = zChain,
+    K = 4, prapivot = mcmc.pars[pivot, , ], p = allocProbs,
+    complete = complete.bivariate.normal.loglikelihood,
```
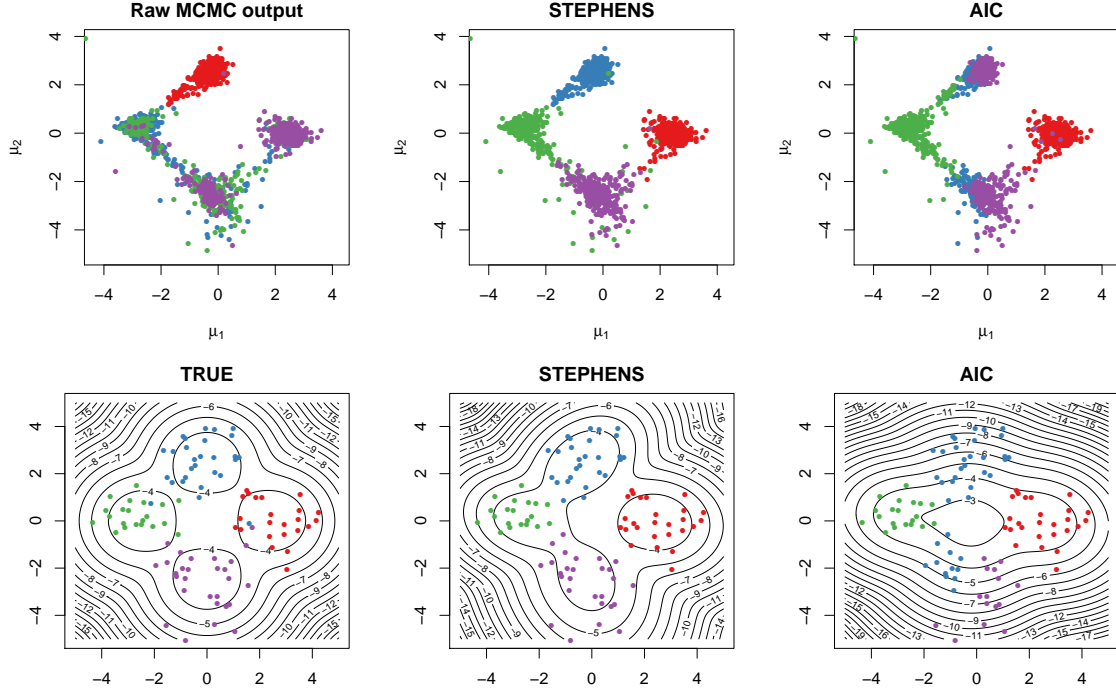
Figure 4: Multivariate dataset 1 ($K = 4$). Top: raw MCMC output of (μ_1, μ_2) and reordered values according to `stephens` and `aic` algorithms for a randomly sampled subset of 500 MCMC iterations. Bottom: True density and clusters of data and corresponding estimates to `stephens` and `aic` algorithms.

```
+     mcmc = mcmc.pars, data = x, sjwinit = pivot, groundTruth = z.real,
+     userPerm = newConstraint$permutations)
```

For the second dataset

```
R> set <- c("STEPHENS", "ECR", "ECR-ITERATIVE-1", "ECR-ITERATIVE-2",
+     "AIC", "DATA-BASED", "USER-PERM")
R> ls <- label.switching(method = set, zpivot = zChain[pivot, ], z = zChain,
+     K = 9, p = allocProbs, mcmc = mcmc.pars, data = x,
+     groundTruth = z.real, userPerm = newConstraint$permutations)
```

The run-times for the reordering part of each algorithm is displayed in Table 3 (last two rows). The simulation study allows to compare all relabeling methods against the ground truth used to generate the data. This is simply done by parsing the option `groundTruth = z.real` to the `label.switching` command (`z.real` corresponds to the true allocations of the observations). Hence, all permutations now are rearranged in order to maximize their similarity with `z.real`. Table 5 displays the coherence between the single best clusterings as returned by `ls$similarity`. Excluding `aic`, there is a strong agreement between the relabeling algorithms and the ground truth, as well as within the relabeling algorithms (note the absolute agreement for the first dataset).

The raw and reordered MCMC output for the means is displayed in Figures 4 and 5 (top). Since most methods produced almost identical results, only `stephens` (for the first dataset)
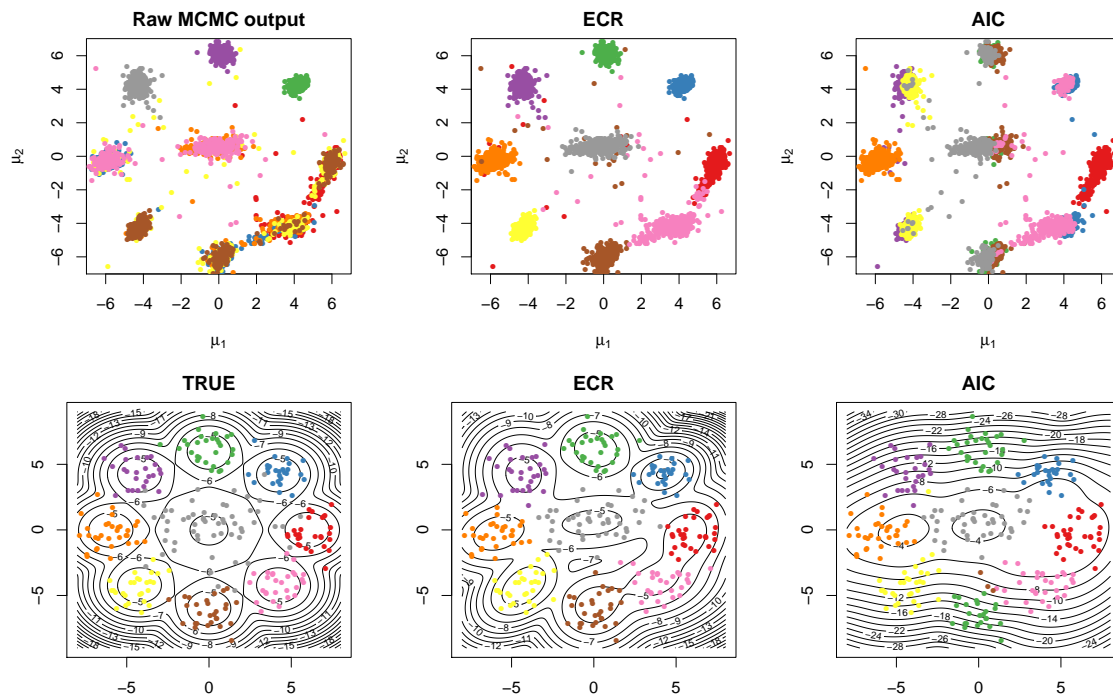
Figure 5: Multivariate dataset 2 ($K = 9$). Top: raw MCMC output of (μ_1, μ_2) and reordered values according to `ecr` and `aic` algorithms for a randomly sampled subset of 500 MCMC iterations. Bottom: True density and clusters of data and corresponding estimates to `ecr` and `aic` algorithms.

and `ecr` (for the second) are shown, along with `aic` which produced different results. The estimated density and single best clustering are displayed in Figures 4 and 5 (bottom). The resulting estimates and single best clusterings reported by `aic` are in stark contrast with the rest of the methods due to the poor performance of the ordering constraint $\mu_{11} < \ldots < \mu_{1K}$. However, a reasonable performance is obtained for the user-defined permutations based on the ordering according to $\mu_{1k} - 2\mu_{2k}$, $k = 1, \ldots, K$, as shown in Table 5.

6. Concluding remarks

The **label.switching** package contains eight relabeling algorithms in order to deal with the problem of non-identifiability in MCMC outputs of mixtures of distributions or hidden Markov models. The input depends on each method, while most of them require information that usually is directly or easily available from the MCMC output. In case that the number of components is small then all algorithms can be applied. When K is large, it is suggested to consider only methods that are optimized using the **lpSolve** routine for the solution of the assignment problem (`dataBased`, `ecr`, `ecr.iterative.1`, `ecr.iterative.2` and `stephens`). Furthermore, all possible simple ordering constraints can be directly applied using the `constraint = "ALL"` argument. In addition, the `"USER-PERM"` option allows the researcher to add new output and make direct comparisons with the available methods.

In practice, the number of components is rarely known. Within a Bayesian framework, K

can be estimated using either Bayes factor approaches (Chib 1995; Carlin and Chib 1995) or trans-dimensional MCMC samplers such as the reversible jump MCMC algorithm of Green (1995) (Richardson and Green 1997; Dellaportas and Papageorgiou 2006; Papastamoulis and Iliopoulos 2009) and the birth-death MCMC sampler of Stephens (2000a). In the first case, a separate MCMC sample for each possible value of K is available, hence each one of them can be directly used as input to the **label.switching** package. In the latter case, the trans-dimensional MCMC sample should be partitioned to subsets for each distinct sampled value of K. Then, in order to make inference conditionally on a given K using the **label.switching** package, the input should correspond to the relevant MCMC draws.

As far as we are concerned, there are no previous efforts for an integrated software of methods dealing with the label switching problem. Given the substantial field of applications of mixture and hidden Markov models as well as the need for making straightforward MCMC inference on complex posterior distributions, the **label.switching** package offers a handy post-processing supplementary tool towards this direction.

Acknowledgments

References

Berkelaar M, others (2015). **lpSolve**: *Interface to* `Lp_solve` *V. 5.5 to Solve Linear/Integer Programs*. R package version 5.6.13, URL https://CRAN.R-project.org/package=lpSolve.

Burkard RE, Dell'Amico M, Martello S (2009). *Assignment Problems*. SIAM e-books. Society for Industrial and Applied Mathematics. doi:10.1137/1.9780898717754.

Cappé O, Moulines E, Ryden T (2005). *Inference in Hidden Markov Models*. Springer Series in Statistics. Springer-Verlag. doi:10.1007/0-387-28982-8.

Carlin B, Chib S (1995). "Bayesian Model Choice via Markov Chain Monte Carlo Methods." *Journal of the Royal Statistical Society B*, **57**(3), 473–484.

Chib S (1995). "Marginal Likelihood from the Gibbs Output." *Journal of the American Statistical Association*, **90**(432), 1313–1321. doi:10.1080/01621459.1995.10476635.

Dellaportas P, Papageorgiou I (2006). "Multivariate Mixtures of Normals with Unknown Number of Components." *Statistics and Computing*, **16**(1), 57–68. doi:10.1007/s11222-006-5338-6.

Dempster JP, Laird NM, Rubin DB (1977). "Maximum Likelihood from Incomplete Data via the EM Algorithm." *Journal of the Royal Statistical Society B*, **39**, 1–38.

Diebolt J, Robert CP (1994). "Estimation of Finite Mixture Distributions through Bayesian Sampling." *Journal of the Royal Statistical Society B*, **39**, 1–37.

Frühwirth-Schnatter S (2001). "Markov Chain Monte Carlo Estimation of Classical and Dynamic Switching and Mixture Models." *Journal of American Statistical Association*, **56**, 363–375. doi:10.1198/016214501750333063.

Frühwirth-Schnatter S (2006). *Finite Mixture and Markov Switching Models: Modeling and Applications to Random Processes.* Springer-Verlag.

Gelfand AE, Smith AFM (1990). "Sampling-Based Approaches to Calculating Marginal Densities." *Journal of American Statistical Association*, **85**, 398–409. doi:10.1080/01621459.1990.10476213.

Genz A, Bretz F, Miwa T, Mi X, Leisch F, Scheipl F, Hothorn T (2016). **mvtnorm**: *Multivariate Normal and t Distributions.* R package version 1.0-4, URL https://CRAN.R-project.org/package=mvtnorm.

Green PJ (1995). "Reversible Jump Markov Chain Monte Carlo Computation and Bayesian Model Determination." *Biometrika*, **82**(4), 711–732. doi:10.1093/biomet/82.4.711.

Grün B (2015). **bayesmix**: *Bayesian Mixture Models with JAGS.* R package version 0.7-4, URL https://CRAN.R-project.org/package=bayesmix.

Jasra A, Holmes CC, Stephens DA (2005). "Markov Chain Monte Carlo Methods and the Label Switching Problem in Bayesian Mixture Modeling." *Statistical Science*, **20**, 50–67. doi:10.1214/088342305000000016.

Leroux B, Putterman ML (1992). "Maximum Penalized Likelihood Estimation for Independent and Markov-Dependent Mixture Models." *Biometrics*, **48**, 545–558. doi:10.2307/2532308.

Marin JM, Mengersen K, Robert CP (2005). "Bayesian Modelling and Inference on Mixtures of Distributions." *Handbook of Statistics*, **25**(1), 577–590. doi:10.1016/s0169-7161(05)25016-2.

Marin JM, Robert CP (2007). *Bayesian Core: A Practical Approach to Computational Bayesian Statistics.* Springer-Verlag. doi:10.1007/978-0-387-38983-7.

McLachlan JG, Peel D (2000). *Finite Mixture Models.* John Wiley & Sons, New York. doi:10.1002/0471721182.

Papastamoulis P (2014). "Handling the Label Switching Problem in Latent Class Models via the ECR Algorithm." *Communications in Statistics – Simulation and Computation*, **43**(4), 913–927. doi:10.1080/03610918.2012.718840.

Papastamoulis P, Iliopoulos G (2009). "Reversible Jump MCMC in Mixtures of Normal Distributions with the Same Component Means." *Computational Statistics & Data Analysis*, **53**(4), 900–911. doi:10.1016/j.csda.2008.10.022.

Papastamoulis P, Iliopoulos G (2010). "An Artificial Allocations Based Solution to the Label Switching Problem in Bayesian Analysis of Mixtures of Distributions." *Journal of Computational and Graphical Statistics*, **19**, 313–331. doi:10.1198/jcgs.2010.09008.

Papastamoulis P, Iliopoulos G (2013). "On the Convergence Rate of Random Permutation Sampler and ECR Algorithm in Missing Data Models." *Methodology and Computing in Applied Probability*, **15**(2), 293–304. `doi:10.1007/s11009-011-9238-7`.

R Core Team (2015). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL `https://www.R-project.org/`.

Redner RA, Walker HF (1984). "Mixture Densities, Maximum Likelihood and the EM Algorithm." *SIAM Review*, **26**, 195–239. `doi:10.1137/1026034`.

Richardson S, Green PJ (1997). "On Bayesian Analysis of Mixtures with an Unknown Number of Components." *Journal of the Royal Statistical Society B*, **59**(4), 731–758. `doi:10.1111/1467-9868.00095`.

Rodriguez CE, Walker SG (2014). "Label Switching in Bayesian Mixture Models: Deterministic Relabelling Strategies." *Journal of Computational and Graphical Statistics*, **23**(1), 25–45. `doi:10.1080/10618600.2012.735624`.

Sperrin M, Jaki T, Wit E (2010). "Probabilistic Relabelling Strategies for the Label Switching Problem in Bayesian Mixture Models." *Statistics and Computing*, **20**(3), 357–366. `doi:10.1007/s11222-009-9129-8`.

Stephens M (2000a). "Bayesian Analysis of Mixture Models with an Unknown Number of Components – An Alternative to Reversible Jump Methods." *The Annals of Statistics*, **28**(1), 40–74. `doi:10.1214/aos/1016120364`.

Stephens M (2000b). "Dealing with Label Switching in Mixture Models." *Journal of the Royal Statistical Society B*, **62**(4), 795–809. `doi:10.1111/1467-9868.00265`.

Tanner M, Wong W (1987). "The Calculation of Posterior Distributions by Data Augmentation." *Journal of the American Statistical Association*, **82**(398), 528–540. `doi:10.1080/01621459.1987.10478458`.

Titterington DM, Smith AFM, Makov UE (1985). *Statistical Analysis of Finite Mixture Distributions*. John Wiley & Sons, Chichester.

Affiliation:

Panagiotis Papastamoulis
Department of Statistics and Insurance Science
University of Piraeus
Greece
E-mail: `papapast@yahoo.gr`