

Bayesian multivariate skew-normal finite mixture model for analysis of infant development trajectories

Carter Allen

Department of Public Health Sciences, Medical University of South Carolina, Charleston, SC, U.S.A

email: allecart@musc.edu

and

Brian Neelon, PhD

Department of Public Health Sciences, Medical University of South Carolina, Charleston, SC, U.S.A

and

Sara Benjamin-Neelon, PhD, MPH, RD

Department of Health, Behavior and Society, Johns Hopkins University, Baltimore, MD, U.S.A

SUMMARY: In studies of infant motor development, a crucial research goal is to identify latent clusters of infants that experience delayed development, as this is a known risk factor for adverse outcomes later in life. However, there are a number of statistical challenges in modeling infant development: the data are typically skewed, exhibit intermittent missingness, and are highly correlated across the repeated measurements collected during infancy. Using data from the Nurture study, a cohort of over 600 mother-infant pairs followed from pregnancy to 12 months postpartum, we develop a flexible Bayesian latent class model for the analysis infant motor development. Our model has a number of attractive features. First, we adopt the multivariate skew normal distribution with cluster-specific parameters that accommodate the inherent correlation and skewness in the data. Second, we model the cluster membership probabilities using a novel Plya-Gamma data-augmentation scheme, thereby improving predictions of the cluster membership allocations. Lastly, we impute missing responses under missing at random assumption by drawing from appropriate conditional skew normal distributions. Bayesian inference is achieved through straightforward Gibbs sampling, and can be carried out in available software such as R. Through simulation studies, we show that the proposed model yields improved inferences over models that ignore skewness. In addition, our imputation method yields improvements compared to conventional missing data methods, including multiple imputation and complete or available case analysis. When applied to Nurture data, we identified two distinct development clusters: one characterized by delayed U-shaped development and a higher percentage of male infants and another characterized by more steady development and a

December 2008

lower percentage of males. The clusters also differed in terms of key demographic variables, such as infant race and maternal pre-pregnancy body mass index. These findings can aid investigators in targeting interventions during this critical early-life developmental window.

KEY WORDS: A key word; But another key word; Still another key word; Yet another key word.

CONTENTS

1 Introduction

1.1 Infant Development Clustering

1.2 Existing Approaches

2 Nurture Study

2.1 Baseline Demographics and Description of Variables

2.2 Statistical Challenges

2.2.1 Skewness of Bayley score residuals

2.2.2 Attrition and Intermittent Missingness

3 Model

3.1 Multivariate Skew Normal Mixture Model

3.2 Multinomial Regression for the Cluster Indicators

3.3 Conditional MSN Imputation

3.4 Bayesian Inference

3.4.1 Prior Specification

3.4.2 Posterior Inference

1. Introduction

1.1 *Infant Development Clustering*

Heterogeneity of treatment effects (HTE) (Lanza and Rhoades, 2013).

1.2 *Existing Approaches*

Mixtures of multivariate non-symmetric distributions such as the multivariate skew-normal (MSN) distribution allow for the nuances of the marginal density to be captured with a more parsimonious set of mixture components. Mixtures of MSN distributions have been dealt with previously in a Bayesian context (Frühwirth-Schnatter & Pyne, 2010), however in these models, focus lies primary on marginal density estimation and inference on the mixture components (i.e. clusters) is not discussed. More recently, the mixtures of skew- t factor analysis (MSTFA) model has been proposed for settings in which cluster-specific inference is of primary interest (Lin *et al.* 2018). However, an important feature not included in the MSTFA is the ability to explain individual-level cluster membership as a function of covariates of interest. Additionally, parameter estimation proposed by Lin *et al.* for the MSTFA relies on a prohibitively complex EM algorithm and does not enjoy the inferential benefits of a Bayesian approach, namely the ability to incorporate prior information into a model and make posterior probability statements. Our proposed model improves on these previous works by estimating parameters in a Bayesian framework as well as including the ability to fit a multinomial logit regression to cluster membership probabilities using a novel application of data augmentation with the Pólya Gamma distribution.

Put lit review of Bayesian PG multinomial logistic regression here

A ubiquitous feature of repeated measures studies is loss of data due to intermittent missingness and attrition. In the Bayesian setting, the standard approach to dealing with missing data is to perform multiple imputation, whereby m imputed data sets are generated from a specified imputation model. After m complete data sets are obtained, parameter

estimates are combined across each data set to produce a final set of parameter estimates (Gelman *et al.* 2013). This approach is not only computationally burdensome, requiring storage and analysis of an $m \times n_{rows} \times n_{cols}$ data array in addition to multiplication of total model run time by a factor of m , but it has been shown to produce unreliable inferences (Zhou and Reiter, 2010). We instead include an “online” imputation step in our Gibbs sampling procedure, whereby missing outcomes are updated at each iteration. This approach greatly increases the number of opportunities for exploration of the missing data parameter space.

2. Nurture Study

2.1 Baseline Demographics and Description of Variables

2.2 Statistical Challenges

2.2.1 Skewness of Bayley score residuals.

2.2.2 Attrition and Intermittent Missingness.

3. Model

3.1 Multivariate Skew Normal Mixture Model

A primary goal of the Nurture study is to identify clusters of infants characterized by distinct motor development trajectories. To address this aim, we propose a flexible finite mixture model that accommodates relevant features of the data, such as skewness and dependence among the responses. To this end, let $\mathbf{y}_i = (y_{i1}, \dots, y_{iJ})^T$ be a $J \times 1$ vector of responses (i.e., Baley scores) for subject i ($i = 1, \dots, n$). For the analysis of the Nurture data, we propose a finite mixture model of the form

$$f(\mathbf{y}_i) = \sum_{k=1}^K \pi_{ik} f(\mathbf{y}_i | \boldsymbol{\theta}_k), \quad (1)$$

where $\boldsymbol{\theta}_k$ is the set of parameters specific to cluster k ($k = 1, \dots, K$) and π_{ik} is a subject-specific mixing weight representing the probability that subject i belongs to cluster k . For now we assume that K is fixed; in Section 4, we discuss model selection strategies for choosing the optimal value of K . We also assume that class membership is fixed throughout the study period, since our focus is to cluster individuals based on their overall developmental patterns over the course of the study. In Section 6, we discuss extensions to allow for class membership to vary over time. **[We could omit these last two sentence – are they really needed? Not sure. Maybe keep for now and think about it.]**

To facilitate posterior inference, we introduce a latent cluster indicator variable z_i taking the value $k \in \{1, \dots, K\}$ with probability π_{ik} . Conditional on $z_i = k$, we assume \mathbf{y}_i is distributed as

$$\mathbf{y}_i | (z_i = k) \sim MSN_J(\boldsymbol{\zeta}_{ki}, \boldsymbol{\alpha}_k, \boldsymbol{\Omega}_k), \quad (2)$$

where $MSN_J(\cdot)$ denotes the J -dimensional multivariate skew normal density, $\boldsymbol{\zeta}_{ki}$ is a $J \times 1$ vector of subject- and cluster-specific location parameters, $\boldsymbol{\alpha}_k$ is a $J \times 1$ vector of cluster-specific skewness parameters, and $\boldsymbol{\Omega}_k$ is a $J \times J$ cluster-specific scale matrix that captures dependence among the J responses. The vector $\boldsymbol{\alpha}_k$ has components α_{kj} , $j = 1, \dots, J$ **[let's**

use kij as our index hierarchy. Please review throughout], that control the skewness of outcome j in cluster k . When $\alpha_k = \mathbf{0}$, the MSN distribution reduces to the multivariate normal distribution $N_J(\zeta_k, \Omega_k)$, where Ω_k is a $J \times J$ covariance matrix.

We can extend model (2) to the regression setting by modeling ζ_{ki} as a function of covariates. Here we adopt a convenient stochastic representation of the MSN density (Azzalini and Dalla Valle, 1996):

$$\mathbf{y}_i | (z_i = k, t_i) = \mathbf{X}_i \boldsymbol{\beta}_k + t_i \boldsymbol{\psi}_k + \boldsymbol{\epsilon}_{ki}, \quad (3)$$

where \mathbf{X}_i is a $J \times Jp$ design matrix that includes potential time-varying covariates (e.g., indicators denoting quarterly visits); $\boldsymbol{\beta}_k = (\beta_{k11}, \dots, \beta_{k1p}, \dots, \beta_{kJ1}, \dots, \beta_{kJp})^T$ is a $Jp \times 1$ vector of cluster- and outcome-specific regression coefficients; $t_i \sim N_{[0, \infty)}(0, 1)$ is a subject-specific standard normal random variable truncated below by zero; $\boldsymbol{\psi}_k = (\psi_{k1}, \dots, \psi_{kJ})^T$ is a $J \times 1$ vector of cluster-specific skewness parameters; and $\boldsymbol{\epsilon}_{ki} \sim N_J(\mathbf{0}, \Sigma_k)$ is a $J \times 1$ vector of error terms. Thus, conditional on t_i and $z_i = k$, \mathbf{y}_i is distributed as $N_J(\mathbf{X}_i \boldsymbol{\beta}_k + t_i \boldsymbol{\psi}_k, \Sigma_k)$. Marginally (after integrating over t_i), \mathbf{y}_i is distributed $MSN_J(\zeta_{ki}, \alpha_k, \Omega_k)$, where through back-transformation

$$\begin{aligned} \zeta_{ki} &= \mathbf{X}_i \boldsymbol{\beta}_k \\ \alpha_k &= \frac{1}{\sqrt{1 - \boldsymbol{\psi}_k^T \Omega_k^{-1} \boldsymbol{\psi}_k}} \boldsymbol{\omega}_k \Omega_k^{-1} \boldsymbol{\psi}_k \quad \text{and} \\ \Omega_k &= \Sigma_k + \boldsymbol{\psi}_k \boldsymbol{\psi}_k^T, \end{aligned}$$

where $\boldsymbol{\omega}_k = \text{Diag}(\sqrt{\omega_{k,11}}, \dots, \sqrt{\omega_{k,JJ}})$ is the $J \times J$ diagonal matrix containing the square root of the diagonal entries of Ω_k . Additional details can be found in Frühwirth-Schnatter and Pyne (2010).

Of note, the MSN density can be expressed more compactly in terms of the matrix skew normal (MatSN) density (Chen and Gupta 2005). As we will see in Section 3.6, the matrix representation of the MSN distribution admits convenient conjugate prior distributions for

the regression parameters and scale matrices, which in turn leads to efficient Gibbs sampling for posterior inference. Let \mathbf{Y}_k be an $n_k \times J$ response matrix with rows \mathbf{y}_i^T , ($i = 1, \dots, n_k$), where $n_k = \sum_{i=1}^n 1_{(z_i=k)}$ is the number of observations in cluster k . From equation (3), it follows that \mathbf{Y}_k is distributed as

$$\begin{aligned}\mathbf{Y}_k &\sim \text{MatSN}_{n_k \times J}(\mathbf{M}_k, \boldsymbol{\alpha}_k, \mathbf{I}_{n_k}, \boldsymbol{\Omega}_k) \\ \text{vec}(\mathbf{M}_k) &= (\boldsymbol{\zeta}_{k1}^T, \dots, \boldsymbol{\zeta}_{kn_k}^T)^T,\end{aligned}$$

where $\boldsymbol{\zeta}_{ki} = \mathbf{X}_i \boldsymbol{\beta}_k$ as in equation (3), $\boldsymbol{\alpha}_k = (\alpha_{k1}, \dots, \alpha_{kJ})^T$, \mathbf{I}_{n_k} is the $n_k \times n_k$ identity matrix, and $\boldsymbol{\Omega}_k$ is the $J \times J$ scale matrix defined above in equation (2). From equation (3), it follows that \mathbf{Y}_k , conditional on the $n_k \times 1$ vector of random effects \mathbf{t}_k , is jointly distributed in matrix form as

$$\mathbf{Y}_k | \mathbf{t}_k \sim \text{MatNorm}_{n_k \times J}(\mathbf{M}_k, \mathbf{I}_{n_k}, \boldsymbol{\Sigma}_k),$$

where $\text{MatNorm}_{n_k \times J}(\cdot)$ denotes a $n_k \times J$ matrix normal density, $\text{vec}(\mathbf{M}_k) = \mathbf{X}_k \boldsymbol{\beta}_k + \mathbf{t}_k \otimes \boldsymbol{\psi}_k$ is an $n_k J \times 1$ mean vector, \mathbf{X}_k is an $n_k J \times Jp$ design matrix, $\boldsymbol{\beta}_k$ is the $(Jp) \times 1$ vector of regression coefficients defined in equation (3), and $\boldsymbol{\Sigma}_k$ is the $J \times J$ conditional covariance of $\boldsymbol{\epsilon}_{ik}$ given in equation (3).

3.2 Multinomial Regression for the Cluster Indicators

To accommodate heterogeneity in the cluster-membership probabilities, we model π_{ik} as a function of covariates using a multinomial logit model

$$\pi_{ik} = \Pr(z_i = k | \mathbf{w}_i) = \frac{e^{\mathbf{w}_i^T \boldsymbol{\delta}_k}}{\sum_{h=1}^K e^{\mathbf{w}_i^T \boldsymbol{\delta}_h}}, \quad k = 1, \dots, K, \quad (4)$$

where \mathbf{w}_i is an $r \times 1$ vector of subject-level covariates, $\boldsymbol{\delta}_k$ is a $r \times 1$ vector of regression parameters associated with membership in cluster k . For identifiability purposes, we fix the reference category $k = K$ and set $\boldsymbol{\delta}_K = \mathbf{0}$. Under this model, $z_i | \mathbf{w}_i \sim \text{Multinom}(1, \boldsymbol{\pi}_i)$, where $\boldsymbol{\pi}_i = (\pi_{i1}, \dots, \pi_{iK})$. During MCMC estimation, the cluster labels z_i are updated from

their multinomial full conditional distribution and used in the remaining MCMC steps as cluster assignments.

3.3 Conditional MSN Imputation

To accommodate missing at random (MAR) responses, we propose a convenient imputation algorithm that can be implemented “online” as part of the Gibbs sampler. In Section 6, we discuss extension to allow for non-ignorable missingness. Suppose \mathbf{y}_i has $q_i \in (1, \dots, J)$ observed values, denoted \mathbf{y}_i^{obs} , and $J - q_i$ intermittent missing values, denoted \mathbf{y}_i^{miss} . We can use of the stochastic representation given in equation (3) to impute \mathbf{y}_i^{miss} from its conditional multivariate normal distribution given (z_i, t_i, Y_i^{obs}) :

$$\begin{aligned} \mathbf{y}_i^{miss} | (z_i = k, t_i, \mathbf{y}_i^{obs}) &\sim N_{J-q}(\boldsymbol{\mu}_{ik}^{miss}, \boldsymbol{\Sigma}_k^{miss}), \text{ where} \\ \boldsymbol{\mu}_{ik}^{miss} &= \\ \boldsymbol{\Sigma}_k^{miss} &= \end{aligned} \tag{5}$$

Carter – work on the above – you will need to define notation and refer readers back to equation (3) as needed. We can discuss next week if needed These results follow from conventional multivariate normal theory. An attractive feature of this imputation algorithm is that it provides more opportunities to explore the parameter space than multiple imputation [**based on summary stats right?**] and avoids multiplicative run-time scaling in m , the number of imputations **Give refs.** In Section 4, we conduct simulation studies to demonstrate that imputing the missing MSN responses improves inferences over complete case analysis.

3.4 Bayesian Inference

3.4.1 Prior Specification. We adopt a fully Bayesian inferential approach and assign prior distributions to all model parameters. Conveniently, all parameters admit conditionally conjugate priors, which greatly improves posterior computation via a data-augmented Gibbs

sampler. For [Give priors for each parameter. Be clear about the conditionally joint prior for β_k and Σ_k . Where appropriate, explain advantages]

3.4.2 Posterior Inference. The above prior specification induces closed-form full conditionals that can be efficiently updated as part of a Gibbs sampler outlined below. Additional details, including derivations can be found in the Web Appendix. [Think about the best way to organize this section. Maybe see my Bayesian Analysis paper for guidance? We can discuss next week.]

Pólya–Gamma Data Augmentation for z_i . The sampler begins by updating the latent cluster indicators z_i ($i = 1, \dots, n$) from its multinomial logit full conditional. To facilitate sampling, we adopt an efficient data-augmentation approach introduced by Polson *et al.* (2013), which expresses the inverse-logit function as a mixture Pólya–Gamma densities. [See my Bayesian Analysis paper for guidance on this part].

[I stopped here]

$$p(\boldsymbol{\delta}_k | \mathbf{Z}, \boldsymbol{\delta}_{k' \neq k}) \propto p(\boldsymbol{\delta}_k) \prod_{i=1}^n \pi_{ik}^{U_{ik}} (1 - \pi_{ik})^{1-U_{ik}}$$

where $p(\boldsymbol{\delta}_k)$ denotes the prior distribution of $\boldsymbol{\delta}_k$, $U_{ik} = 1_{z_i=k}$ is an indicator that subject i belongs to cluster k , and π_{ik} is defined as in Section 3.4. We can rewrite π_{ik} as follows

$$\pi_{ik} = P(U_{ik} = 1) = \frac{e^{\mathbf{w}_i^T \boldsymbol{\delta}_k - c_{ik}}}{1 + e^{\mathbf{w}_i^T \boldsymbol{\delta}_k - c_{ik}}} = \frac{e^{\eta_{ik}}}{1 + e^{\eta_{ik}}}$$

where $c_{ik} = \log \sum_{k' \neq k} e^{\mathbf{w}_i^T \boldsymbol{\delta}_{k'}}$ and $\eta_{ik} = \mathbf{w}_i^T \boldsymbol{\delta}_k - c_{ik}$. We note that the sum $\sum_{k' \neq k} e^{\mathbf{w}_i^T \boldsymbol{\delta}_{k'}}$ includes the reference category, but since we fix $\boldsymbol{\delta}_K = \mathbf{0}$, we have $e^{\mathbf{w}_i^T \boldsymbol{\delta}_K} = 1$, and hence

$$c_{ik} = \log \sum_{k' \neq k} e^{\mathbf{w}_i^T \boldsymbol{\delta}_{k'}} = \log \left(1 + \sum_{k' \notin \{k, K\}} e^{\mathbf{w}_i^T \boldsymbol{\delta}_{k'}} \right)$$

We can use the quantities to re-express the full conditionals for $\boldsymbol{\delta}_k$ as

$$p(\boldsymbol{\delta}_k | \mathbf{Z}, \boldsymbol{\delta}_{k' \neq k}) \propto p(\boldsymbol{\delta}_k) \prod_{i=1}^n \left(\frac{e^{\eta_{ik}}}{1 + e^{\eta_{ik}}} \right)^{U_{ik}} \left(\frac{1}{1 + e^{\eta_{ik}}} \right)^{1-U_{ik}} = p(\boldsymbol{\delta}_k) \prod_{i=1}^n \frac{(e^{\eta_{ik}})^{U_{ik}}}{1 + e^{\eta_{ik}}}$$

which we note is essentially a logistic regression likelihood. We thus apply this Pólya–

Gamma data augmentation scheme to update each δ_k ($k = 1, \dots, K - 1$) one at a time based on the binary indicators U_{ik} .

- Emphasize that PG data augmentation for the multinomial model results in a PG mixture of experts model, which is a computationally efficient way to model edge weights.

Received October 2007. Revised February 2008. Accepted March 2008.